


NATS Community Day




Waldemar Quevedo / [@wallyqs](https://twitter.com/wallyqs)


GopherCon 2017

Agenda

- NATS Protocol
 - Walk through on a very basic client
- Operational aspects from NATS
 - Configuration
 - Clustering
- Monitoring Tools
 - nats-top
 - Prometheus NATS Exporter



Writing a Minimal Client in Go



Implementation (~ 260 lines)

<https://goo.gl/UAJ2Xt>

A Simple Example

Client which connects and publishes a message.

```
package main

import (
    "log"
    "net"
)

func main() {
    // Establish TCP connection.
    conn, err := net.Dial("tcp", "demo.nats.io:4222")
    if err != nil {
        log.Fatalf("Error: %s", err)
    }

    // Publish 'world' message on 'hello' subject.
    pub := []byte("PUB hello 5\r\nworld\r\n")
    _, err = conn.Write(pub)
    if err != nil {
        log.Fatalf("Error: %s", err)
    }

    // Done!
    conn.Close()
}
```

The PUB command

Note: Payload is opaque to the protocol.

```
PUB subject number_of_bytes  
payload
```

When interacting with the server, we have to announce the number of expected number of bytes in the control line:

```
PUB hello 5  
world
```

Though there is a limit in the payload size (by default 1MB). We can get this value from the initial **INFO** encoded in JSON:

```
{
  "auth_required": false,
  "go": "go1.7.4",
  "host": "0.0.0.0",
  "max_payload": 1048576, //
  "port": 4222,
  "server_id": "zrPhBhrjbbUdp2vndDIvE7",
  "tls_required": false,
  "tls_verify": false,
  "version": "0.9.6"
}
```

Sending more bytes than what announced by the server on connect, will cause a client disconnect.

```
NATS @ ~ () $ telnet demo.nats.io 4222
Trying 107.170.221.32...
Connected to demo.nats.io.
Escape character is '^]'.
INFO {"server_id":"zrPhBhrjbbUdp2vndDIvE7","version":"0.9.6","go":"go1.7.4","host":"0.0.0.0","port":4222,
"auth_required":false,"ssl_required":false,"tls_required":false,"tls_verify":false,"max_payload":1048576}
pub hell
```


Processing INFO

```
package main

import (
    "bufio"
    "encoding/json"
    "log"
    "net"
    "strings"
)

func main() {
    // Establishing a connection to a NATS server (blocks!)
    conn, _ := net.Dial("tcp", "demo.nats.io:4222")

    // Read the first line (max control line is 1024 by default btw)
    line, _ := bufio.NewReader(conn).ReadString('\n')

    // Split on first space from control line
    // INFO {...}
    toks := strings.SplitN(line, " ", 2)
    info := struct {
        MaxPayload int `json:"max_payload"`
    }{}
    json.Unmarshal([]byte(toks[1]), &info)

    // => 1048576 bytes by default
    log.Println(info.MaxPayload)
}
```

Registering interest on subject with **SUB**

A **SUB** control line contains the subject and an identifier of the subscription *for the client*.

```
SUB subject identifier
```

e.g. *Register interest on 'foo' and 'hello' subjects using identifiers 1 and 100*

```
SUB foo 1
```

```
SUB hello 100
```

Receiving messages from server

A **MSG** delivered by the server looks like this:

```
MSG subject identifier number_of_bytes  
payload
```

For example, given that we made a subscription on **SUB foo 1** and receiving a message with a '*bar*' payload we would get:

```
MSG foo 1 3  
bar
```

Example consuming messages from a subscription

```
r := bufio.NewReader(conn)
w := bufio.NewWriter(conn)

// Subscribing to a couple of subjects
w.WriteString("SUB foo 1\r\n")
w.WriteString("SUB hello 100\r\n")

// Publishing a couple of commands.
w.WriteString("PUB foo 3\r\nbar\r\n")
w.WriteString("PUB hello 5\r\nworld\r\n")
w.Flush()

for {
    line, err := r.ReadString('\n')
    if err != nil {
        log.Fatalf("Error: %s", err)
    }
    log.Print(line)
}
```

Output would look something like this:

```
INFO {"server_id":"Cxc2FA7Nvx97w0nyoltdRf","version":"0.9.6",...
+OK
+OK
+OK
MSG foo 1 3      # matched: SUB foo 1 and need to read 3 bytes
bar
+OK
MSG hello 100 5  # matched: SUB hello 100 and need to read 5 bytes
world
```

Output would look something like this:

```
INFO {"server_id":"Cxc2FA7Nvx97w0nyoltdRf","version":"0.9.6",...
+OK      # ?
+OK      # ?
+OK      # ?
MSG foo 1 3
bar
+OK      # ?
MSG hello 100 5
world
```

+OK?

These replies from the server are received because by default clients will follow the verbose mode from the protocol, where the server will be acking back with +OK after processing each command.

We can disable this behavior via `CONNECT`.

Disabling verbose mode via CONNECT

```
r := bufio.NewReader(conn)
w := bufio.NewWriter(conn)

// Disable verbose mode
connect := struct {
    Verbose bool    `json:"verbose"`
}{
    Verbose: false,
}
connectOp, _ := json.Marshal(connect)
connectCmd := fmt.Sprintf("CONNECT %s\r\n", connectOp)
w.WriteString(connectCmd)

w.WriteString("SUB foo 1\r\n")
w.WriteString("SUB hello 100\r\n")
w.WriteString("PUB foo 3\r\nbar\r\n")
w.WriteString("PUB hello 5\r\nworld\r\n")
w.Flush()

for {
    line, err := r.ReadString('\n')
    if err != nil {
        log.Fatalf("Error: %s", err)
    }
    log.Print(line)
}
```


Customized protocol interaction after CONNECT

```
INFO {"server_id":"Cxc2FA7Nvx97w0nyoltdRf","version":"0.9.6",...  
MSG foo 1 3  
bar  
MSG hello 100 5  
world
```

Further customizations via CONNECT

We could add a name to the client too...

```
w := bufio.NewWriter(conn)

// Disable verbose mode
connect := struct {
    Name      string `json:"name"`
    Verbose   bool   `json:"verbose"`
}{
    Name:      "gopher",
    Verbose:   false,
}
connectOp, _ := json.Marshal(connect)
connectCmd := fmt.Sprintf("CONNECT %s\r\n", connectOp)
w.WriteString(connectCmd)
w.Flush()
```

Useful to identify more easily a client via nats-top

```
NATS server version 0.9.6 (uptime: 39m23s)
Server:
  Load: CPU:  0.0%  Memory: 9.3M  Slow Consumers: 0
  In:   Msgs: 20  Bytes: 80  Msgs/Sec: 0.0  Bytes/Sec: 0
  Out:  Msgs: 20  Bytes: 80  Msgs/Sec: 0.0  Bytes/Sec: 0

Connections Polled: 1
  HOST          CID  NAME  SUBS  PENDING  MSGS_TO  MSGS_FROM  BYTES_TO  BYTES_FROM
  127.0.0.1:51940 10   gopher 2      0        2         2         8         8
```

At any point in time we may receive from the server any of the following commands:

- MSG
- +OK
- -ERR
- PING
- PONG
- INFO

By default we would get every 2 minutes a **PING** from server, so we must reply back otherwise will be disconnected...


```
INFO {"server_id":"Cxc2FA7Nvx97w0nyoltdRf","version":"0.9.6",...
MSG foo 1 3
bar
MSG hello 100 5
world
PING
PING
-ERR 'Stale Connection'
```

A client connected to NATS is meant to try to always have an established connection to a server in the cluster.

One of the goals from NATS is to provide a *dial tone* for clients to communicate.

Writing a Minimal Client in Go

Protocol Parsing



Parsing the NATS Protocol

Go makes it simple for us to write a naive implementation using `bufio` and `strings`.

```
for {
    line, err := r.ReadString('\n')
    if err != nil {
        log.Fatalf("Error: %s", err)
    }
    args := strings.SplitN(line, " ", 2)
    if len(args) < 1 {
        log.Fatalf("Error: malformed control line")
    }
    op := strings.TrimSpace(args[0])
    switch op {
    case "MSG":
    case "INFO":
    case "PING":
    case "PONG":
    case "+OK":
    case "-ERR":
    default:
        log.Fatalf("Error: malformed control line")
    }
}
```


PING → PONG reply to prevent disconnection

```
for {
    // ...
    switch op {
    case "MSG":
    case "INFO":
    case "PING":
        // Reply back to prevent stale connection error.
        err := w.WriteString("PONG\r\n")
        w.Flush()
        if err != nil {
            log.Fatalf("Error: %s", err)
        }
    case "PONG":
    case "+OK":
    case "-ERR":
    default:
        log.Fatalf("Error: malformed control line")
    }
}
```

Processing INFO

We get one on connect and also when new servers join the cluster.

```
for {
    // ...
    switch op {
    case "MSG":
    case "INFO":
        info := struct {
            MaxPayload int `json:"max_payload"`
            ConnectUrls []string `json:"connect_urls"`
        }{}
        json.Unmarshal([]byte(args[1]), &info)
    case "PING": //
    case "PONG":
    case "+OK":
    case "-ERR":
    default:
        log.Fatalf("Error: malformed control line")
    }
}
```

Handling a received MSG

A MSG may be tagged with an reply inbox, so we have to handle 2 cases:

```
for {
    // ...
    switch op {
    case "MSG":
        // without reply: MSG foo 1 3\r\n
        // with reply:    MSG foo 1 bar 4\r\n
    case "INFO": //
    case "PING": //
    case "PONG":
    case "+OK":
    case "-ERR":
    default:
        log.Fatalf("Error: malformed control line")
    }
}
```

Handling a received MSG

Case 1) Without reply inbox:

```
switch op {
case "MSG":
    var subject, reply string
    var sid, size int

    n := strings.Count(args[1], " ") + 1
    switch n {
    case 3:
        // MSG foo 1 3\r\n
        // bar\r\n
        _, err := fmt.Sscanf(args[1], "%s %d %d",
            &subject, &sid, &size)
        if err != nil {
            log.Fatalf("Error: malformed control line: %s", err)
        }

        // Prepare buffer for the payload with the given size.
        payload := make([]byte, size)
        _, err = io.ReadFull(r, payload)
        if err != nil {
            log.Fatalf("Error: problem gathering bytes: %s", err)
        }
    }
}
```

Handling a received MSG

Case 2) With a reply inbox

```
switch op {
case "MSG":
    var subject, reply string
    var sid, size int

    n := strings.Count(args[1], " ") + 1
    switch n {
    case 4:
        // MSG foo 1 bar 4\r\n
        // quux\r\n
        _, err := fmt.Sscanf(args[1], "%s %d %s %d",
            &subject, &sid, &reply, &size)
        if err != nil {
            log.Fatalf("Error: malformed control line: %s", err)
        }

        // Prepare buffer for the payload with the given size.
        payload := make([]byte, size)
        _, err = io.ReadFull(r, payload)
        if err != nil {
            log.Fatalf("Error: problem gathering bytes: %s", err)
        }
    }
}
```

```
for {  
    // ...  
    switch op {  
    case "MSG": //  
    case "INFO": //  
    case "PING": //  
    case "PONG":  
    case "+OK":  
    case "-ERR":  
    default:  
        log.Fatalf("Error: malformed control line")  
    }  
}
```

Handling PONG

This one is triggered by us when doing a PING, we may disconnect if we have many missing for example.

```
for {
    // ...
    switch op {
    case "MSG": //
    case "INFO": //
    case "PING": //
    case "PONG": // skip for now
    case "+OK":
    case "-ERR":
    default:
        log.Fatalf("Error: malformed control line")
    }
}
```

Handling +OK

Nothing to handle

```
for {
    // ...
    switch op {
    case "MSG": //
    case "INFO": //
    case "PING": //
    case "PONG": //
    case "+OK": // do nothing
    case "-ERR":
    default:
        log.Fatalf("Error: malformed control line")
    }
}
```


Handling -ERR

First argument is the reported error by server

```
for {
    // ...
    switch op {
    case "MSG": //
    case "INFO": //
    case "PING": //
    case "PONG": //
    case "+OK": //
    case "-ERR":
        log.Fatalf("Protocol Error: %s", args[1])
    default:
        log.Fatalf("Error: malformed control line")
    }
}
```

e.g. when not replying to many pings we would get


```
-ERR 'Stale Connection'
```

Mostly working parser

```
for {
    // ...
    switch op {
    case "MSG": //
    case "INFO": //
    case "PING": //
    case "PONG": //
    case "+OK": //
    case "-ERR": //
    default:
        log.Fatalf("Error: malformed control line")
    }
}
```

Writing a Minimal Client in Go

Client API



Public API for the minimal client

```
type Client interface {  
    // Connect establishes connection to NATS.  
    Connect(location string) error  
  
    // Close wraps up connection to NATS.  
    Close()  
  
    // Publish sends a message to a subject.  
    Publish(subj, reply string, payload []byte) error  
  
    // Subscribe registers interest in subject and  
    // delivers messages onto the provided callback.  
    Subscribe(subj, queue string, cb func(subj, reply string, data []byte)) error  
}
```

Connect API

```
func (c *client) Connect(netloc string) error {
    conn, err := net.Dial("tcp", netloc)
    if err != nil {
        return err
    }
    c.conn = conn
    c.r = bufio.NewReader(conn)
    c.w = bufio.NewWriter(conn)
    c.subs = make(map[int]func(string, string, []byte))

    connect := struct {
        Name      string `json:"name"`
        Verbose   bool   `json:"verbose"`
    }{
        Name:      "gopher",
        Verbose:   false,
    }
    connectOp, _ := json.Marshal(connect)
    connectCmd := fmt.Sprintf("CONNECT %s\r\n", connectOp)
    c.w.WriteString(connectCmd)
    c.w.Flush()

    // Spawn goroutine for the parser reading loop.
    go c.runParserLoop()

    return nil
}
```

Basic client which uses the parser

Client which takes a connection and would be called during the parser loop.

```
type client struct {
    conn net.Conn
    r     *bufio.Reader
    w     *bufio.Writer
}

func (c *client) processInfo(info string) { /* TODO */ }

func (c *client) processMsg(subj string, reply string, sid int, payload []byte) {
    /* TODO */
}

func (c *client) processPing() { /* TODO */ }

func (c *client) processPong() { /* TODO */ }

func (c *client) processErr(msg string) { /* TODO */ }
```

Parsing loop → client internal API

```
func (c *client) runParserLoop() {
    for {
        line, _ := c.r.ReadString('\n')
        args := strings.SplitN(line, " ", 2)
        if len(args) < 1 {
            log.Fatalf("Error: malformed control line")
        }
        switch op {
        case "MSG":
            // ...
            c.processMsg(subject, reply, sid, payload)
        case "INFO":
            c.processInfo(args[1])
        case "PING":
            c.processPing()
        case "PONG":
            c.processPong()
        case "+OK":
            // Do nothing
        case "-ERR":
            c.processErr(args[1])
        default:
            log.Fatalf("Error: malformed control line")
        }
    }
}
```

Publish API

```
func (c *client) Publish(subject, reply string, payload []byte) error {
    c.Lock()
    defer c.Unlock()

    pub := fmt.Sprintf("PUB %s %s %d\r\n", subject, reply, len(payload))
    _, err := c.w.WriteString(pub)
    if err != nil {
        return err
    }
    _, err = c.w.Write(payload)
    if err != nil {
        return err
    }
    _, err = c.w.WriteString("\r\n")
    if err != nil {
        return err
    }
    err = c.w.Flush()
    if err != nil {
        return err
    }

    return nil
}
```


Subscribe API / Message Delivery

To support dispatching messages, we need to add a bit more state to the client:

```
type client struct {
    conn net.Conn
    r     *bufio.Reader
    w     *bufio.Writer

    // sid increments monotonically per subscription and it is
    // used to identify a subscription from the client when
    // receiving a message.
    sid int

    // subs maps a subscription identifier to a callback.
    subs map[int]func(subject, reply string, b []byte)

    sync.Mutex
}
```

Subscribe API / Message Delivery

```
func (c *client) Subscribe(subject, queue string,
    cb func(subject, reply string, b []byte),
) error {
    c.Lock()
    defer c.Unlock()
    c.sid += 1
    sid := c.sid

    sub := fmt.Sprintf("SUB %s %s %d\r\n", subject, queue, sid)
    _, err := c.w.WriteString(sub)
    if err != nil {
        return err
    }

    err = c.w.Flush()
    if err != nil {
        return err
    }

    c.subs[sid] = cb

    return nil
}
```

Close API

```
func (c *client) Close() {  
    c.Lock()  
    defer c.Unlock()  
  
    // Send any pending commands to server previous  
    // to closing.  
    c.w.Flush()  
    c.conn.Close()  
}
```

Sample usage of the toy client

Client is roughly 300 lines of code but basic API works!

```
package main

import (
    "log"
    "time"
)

func main() {
    c := &client{}
    err := c.Connect("127.0.0.1:4222")
    if err != nil {
        log.Fatalf("Error: %s", err)
    }
    defer c.Close()
    c.Subscribe("foo", "", func(subject, reply string, data []byte) {
        log.Println("[SUB ] -", subject, reply, "->", string(data))
    })
    for {
        c.Publish("foo", "", []byte("bar"))
        time.Sleep(1 * time.Second)
    }
}
```

Sample usage of the toy client

```
[SUB ] - too -> bar
[MSG ] - subject="hello", reply="", payload=world
[SUB ] - hello -> world
[PONG]
2017-07-04 02:46:58.839492601 -0700 PDT
[MSG ] - subject="foo", reply="", payload=bar
[SUB ] - foo -> bar
[MSG ] - subject="hello", reply="", payload=world
[SUB ] - hello -> world
[PONG]
[PING]
2017-07-04 02:46:59.84080511 -0700 PDT
[MSG ] - subject="foo", reply="", payload=bar
[SUB ] - foo -> bar
[MSG ] - subject="hello", reply="", payload=world
[SUB ] - hello -> world
[PONG]
2017-07-04 02:47:00.841473462 -0700 PDT
[MSG ] - subject="foo", reply="", payload=bar
[SUB ] - foo -> bar
[MSG ] - subject="hello", reply="", payload=world
[SUB ] - hello -> world
[PONG]
[PING]
2017-07-04 02:47:01.842108937 -0700 PDT
[MSG ] - subject="foo", reply="", payload=bar
[SUB ] - foo -> bar
[MSG ] - subject="hello", reply="", payload=world
[SUB ] - hello -> world
[PONG]
```

Operational aspects from NATS

Clustering

Full mesh one hop

```
./go/bin/gnatsd -p 4222 -m 8222 --cluster nats://127.0.0.1:6222 --routes  
nats://127.0.0.1:6222,nats://127.0.0.1:6223,nats://127.0.0.1:6224
```

```
./go/bin/gnatsd -p 4223 -m 8223 --cluster nats://127.0.0.1:6223 --routes  
nats://127.0.0.1:6222,nats://127.0.0.1:6223,nats://127.0.0.1:6224
```

```
./go/bin/gnatsd -p 4224 -m 8224 --cluster nats://127.0.0.1:6224 --routes  
nats://127.0.0.1:6222,nats://127.0.0.1:6223,nats://127.0.0.1:6224
```

Monitoring Tools

nats-top

<https://github.com/nats-io/nats-top/releases/tag/v0.3.2>

Prometheus NATS Exporter

<https://github.com/nats-io/prometheus-nats-exporter/tree/master/walkthrough>